

Contents

Notation	3
Introduction	4
1 Basics of Recursive Bayesian Estimation	5
1.1 Problem Statement	5
1.2 Theoretical solution	6
1.3 Kalman Filter	7
1.4 Particle Filter	9
1.5 Marginalized Particle Filter	11
2 Software analysis	13
2.1 Requirements	13
2.2 Programming paradigms	14
2.2.1 Procedural paradigm	15
2.2.2 Object-oriented paradigm	15
2.2.3 Functional paradigm	16
2.2.4 Other programming language considerations	17
2.3 C++	20
2.4 MATLAB language	21
2.5 Python	23
2.6 Cython	25
2.6.1 Cython Features	26
2.6.2 Performance comparison with C and Python	29
2.6.3 Discussion	31
3 The PyBayes Library	32
3.1 Interpreted and Compiled	32

3.2	Library Layout	32
3.2.1	Random Variable Meta-representation	32
3.2.2	Probability Density Functions	32
3.2.3	Bayesian Filters	32
3.3	Documentation, Testing and Profiling	33
3.4	Performance Comparison with BDM	33
	Conclusion	34
	Bibliography	35

Notation

Throughout this text, following notation is used

\mathbb{N}	set of natural numbers excluding zero
\mathbb{N}_0	set of natural numbers including zero
\mathbb{R}	set of real numbers
t	discrete time moment; $t \in \mathbb{N}_0$
a_t	value of quantity a at time t ; $a_t \in \mathbb{R}^n, n \in \mathbb{N}$
$a_{t t'}$	quantity with two indices: t and t' there is no implicit link between a_t , $a_{t t'}$ and $a_{t'}$
$a_{t:t'}$	sequence of quantities $(a_t, a_{t+1} \dots a_{t'-1}, a_{t'})$
$p(a_t)$	probability density function ¹ of quantity a at time t (unless noted otherwise)
$p(a_t b_{t'})$	conditional probability density function of quantity a at time t given value of quantity b at time t'
$\delta(a)$	Dirac delta function; used exclusively in context of probability density functions to denote discrete distribution within framework of continuous distributions ²
$\mathcal{N}(\mu, \Sigma)$	multivariate normal (Gaussian) probability density function with mean vector μ and covariance matrix Σ

¹for the purpose of this text, probability density function p is multivariate non-negative function $\mathbb{R}^n \rightarrow \mathbb{R}$; $\int_{\text{supp } p} p(x_1, x_2 \dots x_n) dx_1 dx_2 \dots dx_n = 1$

²so that $\int_{-\infty}^{\infty} f(x) \delta(x - \mu) dx = f(\mu)$ and more complex expressions can be derived using integral linearity and Fubini's theorem.

Introduction

TODO motivatin for bayes filtering + a need for a convenient library (rapid prototyping vs. speed)

applications: robotics, navigation, + tracking of toxic plume after radiation accident.

Decision-making being a logical and natural “next step” - beyond the scope of this text.

[proposed citations:[[17](#), [6](#), [8](#), [7](#), [10](#)]]

na co se nezamerujeme: MCMC, Bayes networks

Chapter 1

Basics of Recursive Bayesian Estimation

In following sections the problem of recursive Bayesian estimation (Bayesian filtering) is stated and its analytical solution is derived. Later on, due to practical intractability of the solution in its general form, a few methods that either simplify the problem or approximate the solution are shown.

1.1 Problem Statement

Assume a dynamic system described by a hidden real-valued *state vector* x which evolves at discrete time steps according to a known function f_t (in this text called *process model*) as described by (1.1).

$$x_t = f_t(x_{t-1}, v_{t-1}) \quad (1.1)$$

Variable v_t in (1.1) denotes random *process noise*, which may come from various sources and is often inevitable. Sequence of v_t is assumed to be identically independently distributed random variable sequence.

The state of the system is hidden and can only be observed through a real-valued *observation vector* y that relates to the state x as in (1.2), but adds further *observation noise* w .

$$y_t = h_t(x_t, w_t) \quad (1.2)$$

In (1.2) h_t is known function called *observation model* in this text and w_t is identically independently distributed random variable sequence that denotes observation noise.

The goal of recursive¹ Bayesian estimation is to give an estimate of the state x_t

¹by the word recursive we mean that it is not needed to keep track of the whole batch of previous observations in practical methods, only appropriate quantities from time moments $t - 1$ and t are needed to estimate x_t . However, this does not apply to the derivation of the solution, where the notation of whole batch of observations $y_{1:t}$ is used.

given the observations $y_{1:t}$ provided the knowledge of the functions f_t and h_t . More formally, the goal is to find the probability density function $p(x_t|y_{1:t})$. Theoretical solution to this problem is known and is presented in next section.

1.2 Theoretical solution

At first, we observe that probability density function $p(x_t|x_{t-1})$ can be derived from the process model (1.1) (given the distribution of v_k) and that $p(y_t|x_t)$ can be derived from the observation model (1.2) respectively. (given the distribution of w_k)

Because recursive solution is requested, suppose that $p(x_{t-1}|y_{1:t-1})$ and $p(x_0)$ are known² in order to be able to make the transition $t-1 \rightarrow t$.

In the first stage that can be called *prediction*, *prior* probability density function $p(x_t|y_{1:t-1})$ is calculated without knowledge of y_t . We begin the derivation by performing the reverse of the marginalization over x_{t-1} .

$$p(x_t|y_{1:t-1}) = \int_{-\infty}^{\infty} p(x_t, x_{t-1}|y_{1:t-1}) dx_{t-1}$$

Using chain rule for probability density functions, the element of integration can be split.

$$p(x_t|y_{1:t-1}) = \int_{-\infty}^{\infty} p(x_t|x_{t-1}, y_{1:t-1})p(x_{t-1}|y_{1:t-1}) dx_{t-1}$$

With an assumption that the modelled dynamic system (1.1) possesses *Markov Property*³, $p(x_t|x_{t-1}, y_{1:t-1})$ equals $p(x_t|x_{t-1})$. [1] This leaves us with the result (1.3).

$$p(x_t|y_{1:t-1}) = \int_{-\infty}^{\infty} p(x_t|x_{t-1})p(x_{t-1}|y_{1:t-1}) dx_{t-1} \quad (1.3)$$

As we can see, prior probability density function only depends on previously known functions and therefore can be calculated.

We continue with the second stage that could be named *update*, where new observation y_t is taken into account and *posterior* probability density function $p(x_t|y_{1:t})$ is calculated. Bayes' theorem can be used to derive posterior probability density function (1.4).

$$p(x_t|y_{1:t}) = \frac{p(y_t|x_t, y_{1:t-1})p(x_t|y_{1:t-1})}{p(y_t|y_{1:t-1})} \quad (1.4)$$

According to the observation model (1.2) and assuming Markov property, y_t only depends on x_t . That is $p(y_t|x_t, y_{1:t-1}) = p(y_t|x_t)$. Therefore posterior probability

² $p(x_0)$ can be called initial probability density function of the state vector.

³an assumption of independence that states that system state in time t only depends on system state in $t-1$ (and is not directly affected by previous states).

density function can be further simplified into (1.5).

$$p(x_t|y_{1:t}) = \frac{p(y_t|x_t)p(x_t|y_{1:t-1})}{p(y_t|y_{1:t-1})} \quad (1.5)$$

While both probability density functions in the numerator of (1.5) are already known, $p(y_t|y_{1:t-1})$ found in the denominator can be calculated using the formula (1.6), where marginalization over x_t is preformed. Quantity (1.6) can also be interpreted as *marginal likelihood* (sometimes called *evidence*) of observation. [15]

$$p(y_t|y_{1:t-1}) = \int_{-\infty}^{\infty} p(y_t|x_t)p(x_t|y_{1:t-1}) dx_t \quad (1.6)$$

Computing (1.6) isn't however strictly needed as it does not depend on x_t and serves as a normalising constant in (1.5). Depending on use-case the normalising constant may not be needed at all or may be computed alternatively using the fact that $p(x_t|y_{1:t})$ integrates to 1.

We have shown that so called *optimal Bayesian solution*[1] can be easily analytically inferred using only *chain rule for probability density functions*, *marginalization* and *Bayes' theorem*. (equations (1.3), (1.5) and (1.6) forming the main steps of the solution) On the other hand, using this method directly in practice proves difficult because at least one parametric multidimensional integration has to be performed (in (1.3)), which is (in its general form) hardly tractable for greater than small state vector dimensions.

This is a motivation for various simplifications and approximations among which we have chosen a Kalman filter described in the next section and a family of particle filters described later.

1.3 Kalman Filter

The Kalman filter⁴ poses additional set of strong assumptions on modelled dynamic system, but greatly simplifies the optimal Bayesian solution (1.3), (1.5) into a sequence of algebraic operations with matrices. On the other hand, when these requirements can be fulfilled, there is no better estimator in the Bayesian point of view because the Kalman filter computes $p(x_t|y_{1:t})$ *exactly*.⁵

Assumptions additionally posed on system by the the Kalman filter are:

1. f_t in the process model (1.1) is a linear function of x_t and v_t .
2. $v_t \sim \mathcal{N}(0, Q_t)$ meaning that process noise v_t is normally distributed with zero mean⁶ and with known covariance matrix Q_t .
3. h_t in the observation model (1.2) is a linear function of x_t and w_t .

⁴first presented by Rudolf Emil Kalman in 1960.

⁵not accounting for numeric errors that arise in practical implementations.

⁶zero mean assumption is not strictly needed, it is however common in many implementations.

4. $w_t \sim \mathcal{N}(0, R_t)$ meaning that observation noise w_t is normally distributed with zero mean and with known covariance matrix R_t .
5. initial state probability density function is Gaussian.

It can be proved that if the above assumptions hold, $p(x_t|y_{1:t})$ is Gaussian for all $t > 0$. [11] Furthermore, given assumptions 1. and 2. the process model (1.1) can be reformulated as (1.7), where A_t is real-valued matrix that represents f_t . Using the same idea and assumptions 3. and 4. the observation model (1.2) can be expressed as (1.8), C_t being real-valued matrix representing h_t . Another common requirement used below in the algorithm description is that v_t and w_t are stochastically independent.

$$x_t = A_t x_{t-1} + \hat{v}_{t-1} \quad A_t \in \mathbb{R}^{n,n} \quad n \in \mathbb{N} \quad (1.7)$$

$$y_t = C_t x_t + \hat{w}_t \quad C_t \in \mathbb{R}^{j,n} \quad j \in \mathbb{N} \quad j \leq n \quad (1.8)$$

Note that we have marked the noises v_t and w_t as \hat{v}_t and \hat{w}_t when they are transformed through A_t , respectively C_t matrix. Let also \hat{Q}_t denote the covariance matrix of \hat{v}_t and \hat{R}_t denote the covariance matrix of \hat{w}_t in further text.

At this point we can describe the algorithm of the Kalman filter. As stated above, posterior probability density function is Gaussian and thus can be parametrised by mean vector μ and covariance matrix P . Let us denote posterior mean from previous iteration by $\mu_{t-1|t-1}$ and associated covariance by $P_{t-1|t-1}$ as in (1.9).

$$p(x_{t-1}|y_{1:t-1}) = \mathcal{N}(\mu_{t-1|t-1}, P_{t-1|t-1}) \quad (1.9)$$

Prior probability density function (1.10) can then be calculated as follows: [1]

$$p(x_t|y_{1:t-1}) = \mathcal{N}(\mu_{t|t-1}, P_{t|t-1}) \quad (1.10)$$

$$\mu_{t|t-1} = A_t \mu_{t-1|t-1}$$

$$P_{t|t-1} = A_t P_{t-1|t-1} A_t^T + \hat{Q}_{t-1}$$

Before introducing posterior probability density function it is useful to establish another Gaussian probability density function (1.11) that is not necessarily needed, but is useful because it represents marginal likelihood (1.6).

$$p(y_t|y_{1:t-1}) = \mathcal{N}(\nu_{t|t-1}, S_{t|t-1}) \quad (1.11)$$

$$\nu_{t|t-1} = C_t \mu_{t|t-1}$$

$$S_{t|t-1} = C_t P_{t|t-1} C_t^T + \hat{R}_t$$

The update phase of the Kalman filter can be performed by computing so-called *Kalman gain* matrix (1.12), posterior probability density function (1.13) is then derived from prior one using the Kalman gain K_t and observation y_t . [1]

$$K_t = P_{t|t-1} C_t^T S_{t|t-1}^{-1} \quad (1.12)$$

$$p(x_t|y_{1:t}) = \mathcal{N}(\mu_{t|t}, P_{t|t}) \quad (1.13)$$

$$\mu_{t|t} = \mu_{t|t-1} + K_t (y_t - \nu_{t|t-1})$$

$$P_{t|t} = P_{t|t-1} - K_t C_t P_{t|t-1}$$

In all formulas above A^T denotes a transpose of matrix A and A^{-1} denotes inverse matrix to A . As can be seen, formulas (1.3) and (1.5) have been reduced to tractable algebraic operations, computing inverse matrix⁷ being the most costly one.

It should be further noted that the Kalman filter and described algorithm can be easily enhanced to additionally cope with an *intervention* (or control) vector applied to the system, making it suitable for the theory of decision-making. Numerous generalisations of the Kalman filter exist, for example an *extended Kalman filter* that relaxes the requirement of the linear system by locally approximating a non-linear system with Taylor series. These are out of scope of this text, but provide areas for subsequent consideration.

On the other hand, the assumption of Gaussian posterior probability density function cannot be easily overcome and for systems that show out non-Gaussian distributions of the state vector another approach have to be taken. [1] One such approach can be a Monte Carlo-based *particle filter* presented in the next section.

1.4 Particle Filter

Particle filters represent approximate solution of the problem of the recursive Bayesian estimation, thus can be considered *suboptimal* methods. The underlying algorithm described below is most commonly named *sequential importance sampling (SIS)*. The biggest advantage of the particle filtering is that requirements posed on the modelled system are much weaker than those assumed by optimal methods such as the Kalman filter. Simple form of the particle filter presented in this section (that assumes that modelled system has Markov property) requires only the knowledge of probability density function $p(x_t|x_{t-1})$ representing the process model and the knowledge of $p(y_t|x_t)$ representing the observation model.⁸

The sequential importance sampling approximates posterior density by a weighted empirical probability density function (1.14).

$$p(x_t|y_{1:t}) \approx \sum_{i=1}^N \omega_t^{(i)} \delta(x_t - x_t^{(i)}) \quad (1.14)$$

$$\forall i \in \mathbb{N} \quad i \leq N : \omega_i \geq 0 \quad \sum_{i=1}^N \omega_i = 1$$

In (1.14) $x_t^{(i)}$ denotes value of i -th *particle*: possible state of the system at time t ; $\omega_t^{(i)}$ signifies weight of i -th particle at time t : scalar value proportional to expected

⁷it can be shown that $S_{t|t-1}$ is positive definite given that C_t is full-ranked, therefore the inverse in (1.12) exists.

⁸both probability density functions are generally time-varying and their knowledge for all t is needed, but their representation (parametrised by conditioning variable) is frequently constant in time in practical applications.

probability of the system being in state in small neighbourhood of $x_t^{(i)}$; N denotes total number of particles⁹, a significant tunable parameter of the filter.

As the initial step of the described particle filter, N random particles are sampled from the initial probability density function $p(x_0)$. Let $i \in \mathbb{N}$ $i \leq N$, transition $t - 1 \rightarrow t$ can be performed as follows:

1. for each i compute $x_t^{(i)}$ by random sampling from conditional probability density function $p(x_t|x_{t-1})$ where $x_{t-1}^{(i)}$ substitutes x_{t-1} in condition. This step can be interpreted as a simulation of possible system state developments.
2. for each i compute weight $\omega_t^{(i)}$ using (1.15) by taking observation y_t into account. x_t is substituted by $x_t^{(i)}$ in condition in (1.15). Simulated system states are confronted with reality through observation.

$$\omega_t^{(i)} = p(y_t|x_t)\omega_{t-1}^{(i)} \quad (1.15)$$

3. normalise weights according to (1.16) so that approximation of posterior probability density function integrates to one.

$$\omega_t^{(i)} = \frac{\omega_t^{(i)}}{\sum_{j=1}^N \omega_t^{(j)}} \quad (1.16)$$

Relative computational ease of described algorithm comes with cost: first, the particle filter is in principle non-deterministic because of the random sampling in step 1, in other words, the particle filter is essentially a Monte Carlo method; second, appropriate number of particles N has to be chosen — too small N can lead to significant approximation error while inadequately large N can make the particle filter infeasibly time-consuming. It can be proved that the particle filter converges to true posterior density as N approaches infinity and certain other assumptions hold [4], therefore the number of particles should be chosen as a balance of accuracy and speed.

Only two operations with probability density functions were needed: sampling from $p(x_t|x_{t-1})$ and evaluating $p(y_t|x_t)$ in known point. Sometimes sampling from $p(x_t|x_{t-1})$ is not feasible¹⁰ and/or better results are expected by taking an observation y_t into account during sampling (step 1). This can be achieved by introducing so-called *proposal density* (sometimes *importance density*) $q(x_t|x_{t-1}, y_t)$. Sampling in step 1 then uses $q(x_t|x_{t-1}, y_t)$ instead, where x_{t-1} in condition is substituted by $x_{t-1}^{(i)}$. Weight computation in step 2 have to be replaced with (1.17) that compensates different sampling distribution (every occurrence of x_t , x_{t-1} in the mentioned formula has to be substituted by $x_t^{(i)}$ and $x_{t-1}^{(i)}$ respectively). See [1] for a derivation of these formulas and for a discussion about choosing adequate proposal density.

$$\omega_t^{(i)} = \frac{p(y_t|x_t)p(x_t|x_{t-1})}{q(x_t|x_{t-1}, y_t)}\omega_{t-1}^{(i)} \quad (1.17)$$

⁹ N is assumed to be arbitrary but fixed positive integer for our uses. Variants of the particle filter exist that use adaptive number of particles, these are not discussed here.

¹⁰but can be replaced by evaluation in known point.

Particle filters also suffer from a phenomenon known as *sample impoverishment* or *degeneracy problem*: after a few iterations all but one particles' weight falls close to zero.¹¹ One technique to diminish this is based on careful choice of proposal density (as explained in [1]), a second one is to add additional *resample* step to the above algorithm:

4. for each i resample $x_t^{(i)}$ from approximate posterior probability density function $\sum_{i=1}^N \omega_t^{(i)} \delta(x_t - x_t^{(i)})$ and reset all weights to $\frac{1}{N}$. Given that sampling is truly random and independent this means that each particle is in average copied n_i times, where n_i is roughly proportional to particle weight: $n_i \approx \omega_t^{(i)} N$. Statistics of posterior probability density function are therefore (roughly and on average) maintained while low-weight particles are eliminated.

Step 4 therefore facilitates avoidance of particles with negligible weight by replacing them with more weighted ones. Such enhanced algorithm is known as *sequential importance resampling (SIR)*.

Because particle resampling is computationally expensive operation, a technique can be used where resampling is skipped in some iterations, based on the following idea: a measurement of degeneracy can be obtained by computing an approximate of *effective sample size* N_{eff} at given time t using (1.18). [1]

$$N_{\text{eff}} \approx \left(\sum_{i=1}^N \left(\omega_t^{(i)} \right)^2 \right)^{-1} \quad (1.18)$$

Very small N_{eff} compared to N signifies a substantial loss of “active” particles, which is certainly undesirable as it hurts accuracy while leaving computational demands unchanged. Step 4 is then performed only when N_{eff} falls below certain threshold.

Recursive Bayesian estimation using SIR methods can be applied to a wide range of dynamic systems (even to those where more specialised methods fail) and can be tuned with number of particles N and proposal density q . On the other hand a method specially designed for a given system easily outperforms general particle filter in terms of speed and accuracy.

1.5 Marginalized Particle Filter

Main sources of this section are [12] and [13].

The marginalized particle filter (sometimes *Rao-Blackwellized particle filter*) is an extension to the particle filter that more accurately approximates the optimal Bayesian solution provided that the probability density function representing the process model $p(x_t|x_{t-1})$ can be obtained in a special form. Suppose that the state vector can be divided into two parts a_t and b_t (1.19) and that the process model probability density function can be expressed as a product of two probability density functions (1.20), where $p(a_t|a_{t-1}b_t)$ is analytically tractable (in general). We present

¹¹it has been shown that variance of particle weights continually raises as algorithm progresses. [1]

a simple variant of the marginalized particle filter where, given b_t , process and observation model of the a_t part are linear with normally-distributed noise. The Kalman filter can be used to estimate a_t part of the state vector in this case.

$$x_t = (a_t, b_t) \quad (1.19)$$

$$p(x_t|x_{t-1}) = p(a_t, b_t|a_{t-1}, b_{t-1}) = p(a_t|a_{t-1}, b_t)p(b_t|b_{t-1}) \quad (1.20)$$

The posterior probability density function (1.21) can be represented as a product of a weighted empirical distribution and a normal distribution. Each (i-th) particle is thus associated with its Kalman filter (representing $a_t^{(i)}$ part) and $b_t^{(i)}$ quantity.

$$p(a_t, b_t|y_t) = \sum_{i=1}^N \omega_i p(a_t|y_{1:t}, b_{1:t}^{(i)}) \delta(b_t - b_t^{(i)}) \quad (1.21)$$

In (1.21) N denotes the total number of particles, ω_i denotes weight of i-th particle, $p(a_t|y_{1:t}, b_{1:t}^{(i)})$ is posterior probability density function of i-th Kalman filter and $b_t^{(i)}$ is a value of the b_t part of i-th particle.

The algorithm of the described variant of the marginalized particle filter follows, note the similarities with the ordinary particle filter: at first, generate N b_t random samples from the initial distribution $p(b_0)$. Then the following procedure can be repeated for each measurement y_t :

1. for each i compute $b_t^{(i)}$ by random sampling from conditional probability density function $p(b_t|b_{t-1})$ where $b_{t-1}^{(i)}$ substitutes b_{t-1} in condition.
2. for each i compute the posterior probability density function $p(a_t|y_{1:t}, b_{1:t}^{(i)})$ of the i-th Kalman filter using y_t and $b_t^{(i)}$.
3. for each i update weight ω_i using the formula (1.22) where marginal likelihood of the i-th Kalman filter (1.11) is used.

$$\omega_i = p(y_t|y_{1:t}, b_{1:t}^{(i)}) \omega_i \quad (1.22)$$

4. normalise weights (as described in the previous section).
5. resample particles (as described in the previous section).

It has been demonstrated in various publications that the marginalised particle filter outperforms the ordinary particle filter in both better accuracy and lower computational demands. Where applicable, it therefore forms excellent alternative to traditional particle filtering.

Chapter 2

Software analysis

In this chapter general software development approaches and practices will be confronted with requirements posed on the desired software library for recursive Bayesian estimation. After stating these requirements, feasibility of various programming paradigms applied to our real-world problem is discussed. Continues a comparison of suitable features of 3 chosen programming languages: C++, MATLAB language and Python. Emphasis is put on the Python/Cython combination that was chosen for implementation.

In whole chapter, the term *user* refers to someone (a programmer) who uses the library in order to implement higher-level functionality (such as simulation of dynamic systems).

2.1 Requirements

Our intended audience is a broad scientific community interested in the field of the recursive Bayesian estimation and decision-making. Keeping this in mind and in order to formalise expectations for the desired library for Bayesian filtering, the following set of requirements was developed.

Functionality:

- Framework for working with potentially conditional probability density functions should be implemented including support for basic operations such as product and chain rule. The chain rule implementation should be flexible in a way that for example $p(a_t, b_t | a_{t-1}, b_{t-1}) = p(a_t | a_{t-1}, b_t) p(b_t | b_{t-1})$ product can be represented.
- Basic Bayesian filtering methods such as the Kalman and particle filter have to be present, plus at least one of more specialised algorithms — a marginalized particle filter or non-linear Kalman filter variants.

General:

- Up-to-date, complete and readable API¹ documentation is required. Such docu-

¹Application Programming Interface, a set of rules that define how a particular library is used.

mentation should be well understandable by someone that already understands mathematical background of the particular algorithm.

- High level of interoperability is needed; data input/output should be straightforward as well as using existing solutions for accompanying tasks such as visualising the results.
- The library should be platform-neutral and have to run on major server and workstation platforms, at least on Microsoft Windows and GNU/Linux.
- The library should be Free/Open-source software as it is believed by the authors that such licensing/development model results in software of greatest quality in long term. Framework used by the library should make it easy to adapt and extend the library for various needs.

Usability:

- Initial barriers for installing and setting up the library should be lowest possible. For example a necessity to install third-party libraries from sources is considered infeasible.
- Implementation environment used for the library should allow for high programmer productivity; prototyping new solutions should be quick and cheap (in terms of effort) operation. This requirement effectively biases towards higher-level programming languages.

Performance:

- Computational overhead² should be kept reasonably low.
- Applications built atop of the library should be able to scale well on multi-processor systems. This can be achieved for example by thread-safety of critical library objects or by explicit parallelisation provided by the library.

It is evident that some of the requirements are antagonistic, most prominent example being demand for *low computational overhead* while still offering *high programmer productivity* and rapid prototyping. The task of finding tradeoffs between contradictory tendencies or developing smart solutions that work around traditional limitations is left upon the implementations.

2.2 Programming paradigms

Many programming paradigms exist and each programming language usually suggests a particular paradigm, though many languages let programmers choose from or combine multiple paradigms. This section discusses how well could be three most prominent paradigms (procedural, object-oriented and functional) applied to the software library for Bayesian filtering. Later on additional features of implementation environments such as interpreted vs. compiled approach or argument passing convention are evaluated.

²excess computational costs not directly involved in solving particular problem; for example interpreter overhead.

2.2.1 Procedural paradigm

The procedural paradigm is the traditional approach that appeared along the first high-level programming languages. The procedural programming can be viewed as a structured variant of imperative programming, where programmer specifies steps (in form of orders) needed to reach desired program state. Structured approach that emphasizes dividing the code into logical and self-contained blocks (procedures, modules) is used to make the code more reusable, extensible and modular. Today's most notable procedural languages include C and Fortran.

Most procedural languages are associated with very low overhead (performance of programs compiled using optimising compiler tend to be very close to ideal programs written in assembly code); mentioned languages are also spread and well-known in scientific computing.

On the other hand, while possible, it is considered an elaborate task by the author to write a modular and extensible library in these languages. Another disadvantage is that usually only very basic building blocks are provided by the language — structures like lists and strings have to be supplied by the programmer or a third-party library. This only adds to the fact that the procedural paradigm-oriented languages are commonly not easy to learn and that programmer productivity associated with these languages may be much lower compared to more high-level languages.

2.2.2 Object-oriented paradigm

The object-oriented paradigm extends the procedural approach with the idea of *objects* — structures with procedures (called *methods*) and variables (called *attributes*) bound to them. Other feature frequently offered is *polymorphism* (an extension to language's type system that adds the notion of *subtypes* and a rule that subtype of a given type can be used everywhere where given type can be used) most often facilitated through a concept of *classes*, common models for sets of objects with same behaviour but different payload; objects are then said to be *instances* of classes. A subclass *inherits* methods and attributes from its superclass and can *override* them or add its own. *Encapsulation*, a language mechanism to restrict access to certain object attributes and methods, may be employed by the language to increase robustness by hiding implementation details. In order to be considered object-oriented, statically typed languages (p. 17) should provide *dynamic dispatch*³, an essential complement to polymorphism, for certain or all object methods.

Notable examples of languages that support (although not exclusively) object-oriented paradigm are statically typed C++, Java and dynamically typed (p. 17) MATLAB language, Python, Smalltalk.

Object-oriented features typically have very small overhead compared to procedural code with equal functionality, so additional complexity introduced is the only

³a way of calling methods where the exact method to call is resolved at runtime based on actual (dynamic) object type (in contrast to static object type).

downside, in author’s opinion. We believe that these disadvantages are greatly outweighed by powerful features that object-oriented languages provide (when utilised properly).

It was also determined that the desired library for Bayesian filtering could benefit from many object-oriented techniques: probability density function and its conditional variant could be easily modelled as classes with abstract methods that would represent common operation such as evaluation in a given point or drawing random samples. Classes representing particular probability density functions would then subclass abstract base classes and implement appropriate methods while adding relevant attributes such as border points for uniform distribution. This would allow for example to create generic form of particle filter (p. 9) that would accept any conditional probability density function as a parameter. Bayesian filter itself can be abstracted into a class that would provide a method to compute posterior probability density function from prior one taking observation as a parameter.

2.2.3 Functional paradigm

Fundamental idea of the functional programming is that functions have no side effects — their result does not change or depend on program state, only on supplied parameters. A language where each function has mentioned attribute is called *purely functional* whereas the same adjective is applied to such functions in other languages. This is often accompanied by a principle that all data are immutable (apart from basic list-like container type) and that functions are so-called “first-class citizens” — they can be passed to a function and returned. Placing a restriction of no side-effect on functions allows compiler/interpreter to do various transformations: parallelisation of function calls whose parameters don’t depend on each other’s results, skipping function calls where the result is unused, caching return values for particular parameters.

Among languages specially designed for functional programming are: Haskell, Lisp dialects Scheme and Clojure, Erlang. Python supports many functional programming techniques⁴.

While functional programming is popular subject of academic research, its use is much less widespread compared to procedural and object-oriented paradigms. Additionally, in the author’s opinion, transition to functional programming requires significant change of programmer’s mindset. Combined with the fact that syntax of the mentioned functionally-oriented languages differs significantly from many popular procedural or object-oriented languages, we believe that it would be unsuitable decision for a library that aims for wide adoption.

⁴e.g. functions as first-class citizens, closures, list comprehensions

2.2.4 Other programming language considerations

Apart from recently discussed general approaches to programming, we should note a few other attributes of languages or their implementations that significantly affect software written using them. The first distinction is based on type system of a language — we may divide them into 2 major groups:

statically typed languages

bind object types to *variables*; vast majority of type-checking is done at compile-time. This means that each variable can be assigned only values of given type (subject to polymorphism); most such languages require that variable (function parameter, object attribute) types are properly declared.

dynamically typed languages

bind object types to *values*; vast majority of type-checking is done at runtime. Programmer can assign and reassign objects of arbitrary types to given variable. Variables (and object attributes) are usually declared by assignment.

We consider dynamically typed languages more convenient for programmers — we’re convinced that the possibility of sensible variable reuse and lack of need to declare variable types lets the programmer focus more on the actual task, especially during prototyping stage. This convenience however comes with a cost: dynamic typing imposes inevitable computing overhead as method calls and attribute accesses must be resolved at runtime. Additionally, compiling a program written in statically typed language can reveal many simple programming errors such as calling mistyped methods, even in unreachable code-paths; this is not the case for dynamically-typed languages and we suggest compensating this with more thorough test-suite (code coverage tools can greatly help with creating proper test-suite, see [section 3.3](#) on page 33).

Another related property is interpreted vs. compiled nature; we should emphasize that this property refers to language *implementation*, not directly to the language itself, e.g. C language is commonly regarded as compiled one, several C interpreters however exist. We use the term “language is compiled/interpreted” to denote that principal implementation of that language is compiled, respectively interpreted.

compiled implementations

translate source code directly into machine code suitable for given target processor. Their advantage is zero interpreter overhead. Developers are required to install a compiler (and perhaps a build system) or an IDE⁵ used by given project (library) to be able to modify it. Write-build-run-debug cycle is usually longer in comparison to interpreted implementations.

interpreted implementations

either directly execute commands in source code or, more frequently, translate source code into platform-independent *intermediate representation* which is afterwards executed in a *virtual machine*. We may allow the translate and execute steps to be separated so that Java and similar languages can be included. Advantages include usually shorter write-run-debug cycle that speeds up de-

⁵Integrated Development Environment

velopment and portable distribution options. Interpreted languages have been historically associated with considerable processing overhead, but *just-in-time compilation*⁶ along with *adaptive optimisation*⁷ present in modern interpreters can minimise or even reverse interpreter overhead: Paul Buchheit have shown⁸ that second and onward iterations of fractal-generating Java program were actually 5% faster than equivalent C program. We have reproduced the test with following results: Java program was 10% slower (for second and subsequent iterations) than C program and 1600% slower when just-in-time compilation was disabled. Complete test environment along with instructions how to reproduce it be found in the [examples/benchmark_c_java](#) directory in the PyBayes source code repository.

There exists a historic link between statically typed and compiled languages, respectively dynamically typed and interpreted languages. Java which is itself statically typed and it's major implementation is interpreted and Erlang's (which is dynamically typed) compiled HiPE⁹ implementation are some examples of languages that break the rule. We believe that this historic link is the source of a common misconception that interpreted languages are inherently slow. Our findings (see also Python/Cython/C benchmark on p. 29) indicate that the source of heavy overhead is likely to be the dynamic type system rather than overhead of modern just-in-time interpreters. In accordance with these findings, we may conclude that choice of language implementation type should rather be based on development and distribution convenience than on expected performance.

Each programming language may support one or more following function call conventions that determine how function parameters are passed:

call-by-value convention

ensures that called function does not change variables passed as parameters from calling function by copying them at function call time. This provides clear semantics but incurs computational and memory overhead, especially when large data structures are used as parameters. As a form of optimisation, some language implementations may employ copy-on-write technique so that variables are copied only when they are mutated from within called function, thus saving space and time when some parameters are only read from.

call-by-reference convention

hands fully-privileged references to parameters to called function. These references can be used to modify or assign to parameters within called function and these changes are visible to calling function. This approach minimises function call overhead but may appear confusing to a programmer when local variable is changed "behind her back" unexpectedly. On the other hand, call-by-reference allows for programming techniques impossible with call-by-value alone (e.g. a

⁶interpreter feature that translates portions of bytecode into machine code at runtime.

⁷a technique to use profiling data from recent past (collected perhaps when relevant portion of code was run in interpreted mode) to optimise just-in-time compiled code.

⁸<http://paulbuchheit.blogspot.com/2007/06/java-is-faster-than-c.html>

⁹The High-Performance Erlang Project: <http://www.it.uu.se/research/group/hipe/>

function that swaps two values).

call-by-object (call-by-sharing) convention

can be viewed as a compromise between call-by-value and call-by-reference: parameters are passed as references that can be used to modify referred objects (unless marked immutable), but cannot be used to assign to referred objects (or this assignment is invisible to calling function). When an object is marked as immutable, passing this object behaves like call-by-value call without copying overhead (in the calling function point of view). Java and Python use call-by-object as their sole function calling method¹⁰ and both mark certain elementary types (most prominently numbers and strings) as immutable. C's pointer-to-const and C++'s reference-to-const parameters can be viewed as call-by-object methods where referred objects are marked as immutable in called function scope.

We suggest that a language that supports at least one of call-by-reference or call-by-object conventions is used for the desired recursive Bayesian estimation library; while call-by-value-only languages can be simpler to implement, we are convinced that they impose unnecessary restrictions on the library design and cause overhead in places where it could be avoided.

Last discussed aspect of programming languages relates to memory management:

garbage-collected languages

provide memory management in the language itself. This fact considerably simplifies programming as programmer doesn't need to reclaim unused memory resources herself. Another advantage is that automatic memory management prevents most occurrences of several programming errors: memory leaks,¹¹ dangling pointers¹² and double-frees.¹³ Two major approaches to garbage collection exist and both incur runtime computational or memory overhead. *Tracing garbage collector* repeatedly scans program heap¹⁴ memory for objects with no references to them, then reclaims memory used by these objects. Program performance may be substantially impacted while tracings garbage collector performs its scan; furthermore the moment when garbage collector fires may be unpredictable. *Reference counting* memory management works by embedding an attribute, *reference count*, to each object that could be allocated on heap and then using this attribute to track number of references to given object. When reference count falls to zero, the object can be destroyed. Reference counting adds small memory overhead per each object allocated and potentially significant computational overhead as reference counts have to be kept up-to-date. However, techniques exist that minimise this overhead, for example those mentioned in [9].

¹⁰python case: <http://effbot.org/zone/call-by-object.htm>

¹¹an error condition when a region of memory is no longer used, but not reclaimed.

¹²a pointer to an object that has been already destroyed; such pointers are highly error-prone.

¹³an error condition where a single region of memory is reclaimed twice; memory corruption frequently occurs in this case.

¹⁴an area of memory used for dynamic memory allocation.

non garbage-collected languages

put the burden of memory management on shoulders of the programmer: she is responsible for correctly reclaiming resources when they are no longer in use. The advantages are clear: no overhead due to memory management, probably also smaller complexity of language implementation. However, as mentioned earlier, languages without automatic memory management make certain classes of programmer errors more likely to occur.

In our view, convenience of garbage-collected languages outweighs overhead they bring for a project like a library for recursive Bayesian estimation targeting wide adoption. We also believe that automatic memory management can simplify library design and its usage as there is no need to specify who is responsible for destroying involved objects on the library side and no need to think about it at the user side.

2.3 C++

C++ is regarded as one of the most popular programming languages today, along with Java and C;¹⁵ it combines properties of both low-level and high-level languages, sometimes being described as intermediate-level language. C++ extensively supports both procedural and class-based object-oriented paradigm, forming a multi-paradigm language; generic programming is implemented by means of *templates*, which allow classes and functions to operate on arbitrary data types while still being type-safe. C++ is statically-typed, all major implementations are compiled, supports call-by-value (the default), call-by-reference and a variant of call-by-object function call conventions. C++ lacks implicit garbage collection for heap-allocated data — the programmer must reclaim memory used by those objects manually; use of *smart pointers*¹⁶ may although help with this task. C++ is almost 100% compatible with the C language in a way that most C programs compile and run fine then compiled as C++ programs. C++ also makes it easy to use C libraries without a need to recompile them. [16]

When used as an implementation language for the desired library for recursive Bayesian estimation, we have identified potential advantages of the C++ language:

low overhead

C++ was designed to incur minimal overhead possible. In all benchmarks we've seen (e.g. The Computer Language Benchmarks Game¹⁷), it is hard to outperform C++ by a significant margin (Fortran and assembly code would be candidates for that).

widespread

C/C++ code forms large part of the software ecosystem. Thanks to that, in-

¹⁵TIOBE Programming Community Index for July 2011: <http://www.tiobe.com/index.php/content/paperinfo/tpci/index.html>

¹⁶a template class that behaves like a pointer through use of operator overloading but adds additional memory management features such as reference counting

¹⁷<http://shootout.alioth.debian.org/>

credible number of both proprietary and free IDEs, debuggers, profilers and other related coding tools is available. This fact makes development more convenient.

libraries

Thanks to C++ popularity, several high-quality libraries for numerical calculations/computer algebra are available, many of them are free software or free to use. These are for example C interfaces to BLAS¹⁸ and LAPACK¹⁹ (both low-level and fixed function), higher-level IT++²⁰ built atop of BLAS/LAPACK or independent template-based library Eigen.²¹ Additionally, OpenMP²² can be used to parallelise existing algorithms without rewriting them.

However, using C++ would, in our opinion, bring following major drawbacks:

diversity

While there are many C/C++ libraries for specific tasks (such as data visualisation), it may prove difficult in our opinion to combine them freely as there are no *de facto* standard data types for e.g. vectors and matrices — many libraries use their own.

learning curve

C++ takes longer to learn and even when mastered, programmer productivity is subjectively lower compared to very high-level languages. We also fear that many members of our intended audience are simply unwilling to learn or use C++.

Moreover, discussion about statically-typed, compiled and non-garbage-collected languages from previous section also apply. Due to this, we have decided not to use C++ if an alternative with reasonable overhead is found.

Several object-oriented C++ libraries for recursive Bayesian estimation exist: Bayes++²³, BDM [19] and BFL [5]. BDM library is later used to compare performance of Cython, C++ and MATLAB implementations of the Kalman filter, see [section 3.4](#) on page 33.

2.4 MATLAB language

MATLAB language is a very high-level language used exclusively by the MATLAB²⁴ environment, a proprietary platform developed by MathWorks.²⁵ MATLAB language extensively supports procedural programming paradigm and since version 7.6

¹⁸Basic Linear Algebra Subprograms: <http://www.netlib.org/blas/>

¹⁹Linear Algebra PACKage: <http://www.netlib.org/lapack/>

²⁰<http://itpp.sourceforge.net/>

²¹<http://eigen.tuxfamily.org/>

²²Open Multi-Processing API and libraries: <http://openmp.org/>

²³<http://bayesclasses.sourceforge.net/>

²⁴<http://www.mathworks.com/products/matlab/>

²⁵<http://www.mathworks.com/>

(R2008a) class-based object oriented paradigm is also fully supported.²⁶ MATLAB language is dynamically-typed, interpreted language with automatic memory management.

MATLAB language possesses, in our belief, following favourable attributes when used to implement the desired library for Bayesian filtering:

popularity among academia

While MATLAB language is not as widespread as C++ on the global scale, it is very popular in scientific community, our intended audience.

performance

MATLAB language is very well optimised for numerical computing.

wide range of extensions

High number of well integrated extension modules (toolboxes) is bundled with MATLAB or available from third parties. This makes associated tasks such as data visualisation particularly straightforward.

rapid development

Being a very high-level language, we expect programmer productivity in the MATLAB language being fairly high. MATLAB environment is itself a good IDE and its interactive shell fosters rapid prototyping.

Following disadvantages of the MATLAB language were identified:

vendor lock-in

MATLAB is commercial software; free alternatives such as GNU Octave²⁷, Scilab²⁸ or FreeMat²⁹ exist, however all of them provide only limited compatibility with the MATLAB language. Developing for a non-standard proprietary platform always imposes risks of the vendor changing license or pricing policy etc.

problematic object model

We have identified in [subsection 2.2.2](#) that object-oriented approach is important for a well-designed and usable library for Bayesian filtering. Nonetheless MATLAB's implementation of object-oriented programming is viewed as problematic by many, including us. For example, function call parameter passing convention is determined by the object class/data type — MATLAB distinguishes *value classes* that have call-by-value semantics and *handle classes* that have call-by-object semantics.³⁰ The resulting effect is that calling identical function with otherwise equivalent value and handle classes can yield very different behaviour.

hard-coded call-by-value semantics

2D array, a very central data-type of the MATLAB language, has call-by-value function call convention hard-coded; this results in potentially substantial function call overhead. Although current MATLAB versions try to minimise copy-

²⁶<http://www.mathworks.com/products/matlab/whatsnew.html>

²⁷<http://www.gnu.org/software/octave/>

²⁸<http://www.scilab.org/>

²⁹<http://freemat.sourceforge.net/>

³⁰call-by-object semantics tested in version 7.11 (R2010b).

ing by employing copy-on-write technique³¹ or performing some operations in-place,³² our tests have shown that even combining these techniques doesn't eliminate unnecessary copying overhead which we believe is the main source of grave performance regression of object-oriented code with regards to imperative code; see [section 3.4](#) on page 33.

We consider presented drawbacks significant and therefore decided not to use the MATLAB language for the desired Bayesian filtering library. BDM library [19] contains both object oriented and imperative implementation of the Kalman filter in the MATLAB language; these are compared with our implementation in [section 3.4](#).

2.5 Python

Python³³ is a very high level programming language designed for outstanding code readability and high programmer productivity actively developed by the Python Software Foundation.³⁴ Python extensively supports procedural and class-based object-oriented programming paradigms and some features of the functional programming. Python is dynamically-typed language with automatic memory management that exclusively employs call-by-object function call parameter passing convention; elementary numeric types, strings and tuples are immutable³⁵ so that this approach doesn't become inconvenient.

Principal Python implementation, CPython, is written in C, is cross-platform and of interpreted type: it translates Python code into bytecode which is subsequently executed in a virtual machine. Many alternative implementations are available, to name a few: Jython³⁶ that translates Python code into Java bytecode (itself written in Java), IronPython³⁷ itself implemented on top of the .NET Framework, just-in-time compiling PyPy³⁸ written in Python itself or Cython which is described in greater detail in the next section. All the mentioned implementations qualify as free/open-source software.

Python language is bundled with a comprehensive standard library so that writing new projects is quick from the beginning. Two major Python versions exists: Python 2, considered legacy and receiving only bugfix updates, and Python 3, actively developed and endorsed version that brings a few incompatible changes to the language syntax and to the standard library. Porting Python 2 code to version 3 is however usually straightforward and can be automated to a great extent with tools bundled with Python 3.

³¹<http://blogs.mathworks.com/loren/2006/05/10/memory-management-for-functions-and-variables/>

³²<http://blogs.mathworks.com/loren/2007/03/22/in-place-operations-on-data/>

³³<http://www.python.org/>

³⁴<http://www.python.org/psf/>

³⁵<http://docs.python.org/reference/datamodel.html>

³⁶<http://www.jython.org/>

³⁷<http://ironpython.net/>

³⁸<http://pypy.org/>

In our belief, Python shows following favourable attributes when used for the desired Bayesian filtering library:

development convenience, readability, rapid prototyping

Python developers claim that Python is an easy to learn, powerful programming language and our experience confirms their claims. Python code is easy to prototype, understand and modify in our opinion; prototyping is with bundled interactive Python shell. While all these statements are subjective, they are shared among many.³⁹ For example a statement $x \leq y \leq z$ has its mathematical meaning, which is unusual for programming languages.

NumPy, SciPy, Matplotlib

NumPy project⁴⁰ is the de facto standard Python library for numeric computing; NumPy provides N-dimensional array type that is massively supported in very high number of projects. Parts of NumPy are written in C and Cython for speed. SciPy⁴¹ extends NumPy with more numerical routines. Matplotlib⁴² is powerful plotting library that natively supports SVG output. Combining these three and Python gives a very vital MATLAB alternative.

interoperability with C

CPython makes it possible to write modules⁴³ in C⁴⁴ (that are then called *extension modules*). Cython makes it easy and convenient to write extension modules. Sadly, alternative implementations PyPy, Jython and IronPython don't currently fully support extension modules and are therefore ruled-out for our purposes because they in turn don't support NumPy.⁴⁵

interoperability with MATLAB

SciPy contains procedures to load and save data in MATLAB .mat format; Matplotlib includes programming interface that resembles MATLAB's plotting procedures.

On the other hand, a few downsides exist:

overhead

CPython implementation shows significant computational overhead, especially for numerical computations; CPython doesn't currently utilise any form of just-in-time compiling. NumPy is often used to trade off computational overhead for memory overhead: The Computer Language Benchmarks Game⁴⁶ contains an example where a program heavily using NumPy is $20\times$ faster but consumes $12\times$ more memory than a program that performed the same task and used solely the Python standard library. We have reproduced the benchmark with similar

³⁹<http://python.org/about/quotes/>

⁴⁰<http://www.numpy.org/>

⁴¹<http://www.scipy.org/>

⁴²<http://matplotlib.sourceforge.net/>

⁴³module in python sense is a code unit with its own namespace, normally each module corresponds to a .py file.

⁴⁴<http://docs.python.org/extending/index.html>

⁴⁵PyPy and IronPython are nonetheless interesting for future consideration as both have NumPy support actively worked on.

⁴⁶<http://shootout.alioth.debian.org/>

results; mentioned programs can be found in the [examples/benchmark_py_numpy](#) directory in the PyBayes source code repository. We still consider this workaround suboptimal.

peculiar parallelisation

While Python natively supports threads and they are useful for tasks such as background I/O operations, Python threads don't scale on multiprocessor systems for CPU-bound processing; such code often runs at single-processor speed (when run in CPython). The reason behind that is that CPython employs a global-interpreter-lock (GIL) to assure that only one thread executes Python bytecode at a time.⁴⁷ This restriction can be worked around by using multiple python interpreters that communicate with each other; Python module *multiprocessing* makes it almost as convenient as using threads.

Python compares favourably to other implementation environments presented before in our view; sole major obstacle being excessive overhead of the CPython interpreter. It is discussed how this issue can be solved using Cython in the next section.

We haven't found any Python library for recursive Bayesian estimation that would fulfil the requirements presented at the beginning of this chapter.

2.6 Cython

Cython [2] is both an extension to the Python language and an implementation of it (a compiler). Cython works by translating Cython modules (.pyx or .py files) into the C language which is then compiled to form binary Python extension modules (dynamically loaded libraries — .dll files on Windows and .so files on UNIX platforms). Cython aims to be a strict superset of Python in a way that a valid Python module is also a valid Cython module that behaves equally. Current development snapshot of Cython virtually achieves this goal as it successfully passes 97.6% of the Python 2.7.1 regression test suite. Additionally, interpreted Python modules and Cython-compiled modules can be mixed and interchanged freely as Cython-compiled code can transparently interact with interpreted Python code and vice-versa. However, Cython is not a replacement of CPython — Cython-compiled modules need to be executed by CPython because they make heavy use of CPython internals written in C (C code emitted by the Cython compiler is largely composed of calls to functions from CPython C API); Cython-compiled modules merely circumvent the virtual machine (bytecode interpreter) part of CPython, thus virtually eliminating interpreter overhead. This is in fact probably the most serious limitation of Cython — it is tightly bound to one specific (although the most prevalent) Python implementation, CPython.

Another key feature Cython supports is static typing of variables, function parameters and object attributes. In addition, Cython allows statically typed variables to be native C types such as `int`, `char *` or `struct sockaddr` in addition to Python types;

⁴⁷<http://docs.python.org/glossary.html>

automatic conversion on Python/C boundary is provided by Cython for C types with Python equivalents, C `double` is wrapped as Python `float` and C `long` is wrapped as Python `int` for example. The main purpose of static typing are impressive speed gains; an extreme case exploited by our test case showed $65\times$ speed increase of typed Cython code compared to untyped Cython code. In important thing to mention is that static typing is completely *voluntary* for the programmer; recent Cython versions also support experimental basic type inference⁴⁸ (that is guaranteed not to change semantics). Simply put, static typing prevents significant overhead caused by highly dynamic nature of Python.

2.6.1 Cython Features

We continue by a brief technical description of some Cython extensions to Python and other features so that Cython's benefits can be evaluated as a whole later:

cdef variables

Cython supports so-called *cdef variables* in both global (module) scope and function (method) scope that can be typed; in addition to Python types, they can also have native C types. Cdef variables aren't visible to Python code, but some overhead is eliminated as they are, for example, not reference-counted.

cdef and cpdef functions

In addition to traditional Python functions that are defined in code as `def function(x)` and use rather expensive calling convention (e.g. all their positional arguments are packed into a tuple and all their keyword arguments are packed into a dictionary on each call) Cython supports so-called *cdef functions* and *cpdef functions*. Cdef functions (defined as `cdef function(x)`, hence the name) use native C calling convention and are only callable from Cython code, with greatly reduced overhead; their return value can be typed (parameters can be typed even for traditional Python functions). Cpdef functions (defined as `cpdef function(x)`) are same as cdef functions but in addition a Python wrapper around the C function with the same name is generated; such function is callable from Python (with overhead) and fast to call from Cython, cpdef functions therefore combine the benefits of both. Neither cdef or cpdef functions can be Python class methods, but see the next entry.

extension types (cdef classes)

In addition to traditional Python classes (defined as `class ClassName(SuperClass)`) that store their attributes and methods in a class dictionary, which leads to inefficient attribute lookups and method calls, Cython supports so-called *cdef classes* (or extension types) that are defined as `cdef class ClassName(SuperClass)` and use C structs to store their attributes and method table, a significantly faster approach than Python dictionary. Cdef classes can also have cdef and cpdef methods as their members; they can also contain cdef variables as attributes with optional additional modifiers `public` (read-write from Python) or

⁴⁸a way to guess and prove type of certain variable using code analysis

`readonly` (read-only from Python). `Cdef` classes are, in contrast to similarly named `cdef` functions, visible to Python code where they appear as built-in types. Compared to traditional Python classes, `cdef` classes are however subject to 2 major limitations: all their attributes have to be statically declared (as opposed to dynamically created at runtime) plus multiple inheritance cannot be used. One can overcome both these limitations by subclassing them in Python, which is indeed possible.

interfacing C/C++ code

As described in [3], Cython can be used to conveniently call C or C++⁴⁹ library functions. Suppose we want to call the `sin()` function from the C standard library that is defined in `math.h`. Following Cython module accomplishes that:

```
sin_wrapper.pyx
cdef extern from math.h:
    double sin(double x)

sin(0.123) # call C function directly from Cython

cpdef double sin_wrapper(double x):
    return sin(x)
```

In the example above, `sin_wrapper()` can be called from Python code. Cython is therefore an excellent tool to create Python wrappers around C/C++ libraries or use them directly.

NumPy support

As was noted above, many core parts of NumPy are written in C (or Cython), including the crucial data-type, N-dimensional array. Cython provides explicit support for NumPy data-types and core functions allowing the programmer to take advantage of the speed gains of both. [14]

pure Python mode

It is clear that in order to make use of all important features of Cython, one has to use syntax that is no longer valid Python code; this can be disadvantageous in many cases. To combat this, Cython offers so-called *pure Python mode* where Cython-specific syntax is wrapped in Python constructs, e.g. `cdef int i` can be alternatively formulated as `x = cython.declare(cython.int)` — such constructs are implemented by Cython shadow code to be no-ops when interpreted by Python and treated accordingly when compiled by Cython.

Another variant of pure Python mode is to use so-called *augmenting files*, `.pxd` files that accompany equally-named `.py` files. Such augmenting file can contain declarations of variables, functions and classes that appear in associated `.py` files and add Cython-specific features to them. For example, suppose that it is desired to have a function that is inexpensive to call from Cython but still providing full Python compatibility, one can write following Python module:

⁴⁹C code emitted by the Cython compiler is compilable also when treated as C++, thus allows interfacing with C++ code.

```

module.py
def f(x):
    return x*x

```

And associated augmenting file:

```

module.pxd
cdef double f(double x)

```

It is forbidden to provide implementation in .pxd files. This alternative approach is advantageous to the first one because it doesn't require Cython to be installed on a target machine (on the other hand, some very special Cython features currently cannot be expressed in augmenting files). .pxd files also serve for sharing declarations between Cython modules in a way similar to C header (.h) files.⁵⁰

parallelisation

One recent feature of current development snapshots of Cython is a native support for C-level parallelisation through OpenMP⁵¹ — Cython adds new `prange()` function that is similar to the Python `range()` function except that loop statements are executed in parallel. Currently only loops whose statements can be called without holding the GIL⁵² can be efficiently parallelised. Our tests have shown that such parallel loops scale well with the number of processors, see [subsection 2.6.2](#) on page 29.

Python 3 compatibility

Cython provides full compatibility with Python 3 in particularly robust way: Cython modules can be written in both Python 2 or Python 3 and C files that Cython produces can be compiled against both CPython 2 or CPython 3, independently from the source file version.⁵³ This effectively makes Cython a 2-way compatibility bridge between Python 2 and Python 3.

There are many areas where Cython can get better, however most of them are just missed optimisation possibilities. We mention 2 major cases where Cython currently doesn't accept valid Python constructs: the *yield* statement⁵⁴ used to create generator functions is currently unsupported and *generator expressions*⁵⁵ are supported only in special cases. Other minor inconveniences exist but we don't consider them worth discussing here. Another important attribute of Cython is, in our belief, its active developer community, for example we've discovered a bug related to Cython's pure Python mode capability that was fixed upon providing a test-case.⁵⁶

⁵⁰http://docs.cython.org/src/userguide/sharing_declarations.html

⁵¹The OpenMP API specification for parallel programming: <http://openmp.org/>

⁵²global-interpreter lock <http://docs.python.org/glossary.html>

⁵³<http://wiki.cython.org/FAQ>

⁵⁴http://docs.python.org/reference/simple_stmts.html

⁵⁵<http://docs.python.org/reference/expressions.html>

⁵⁶http://trac.cython.org/cython_trac/ticket/583

2.6.2 Performance comparison with C and Python

In order to evaluate performance of Cython, we've conducted a benchmark where equivalent Python, Cython-compiled Python, Cython and C programs are compared. Inspired by the Cython tutorial, the test program performs simple numerical integration of the function x^2 from 0 to 3. Python version of the test program is shown here, complete test environment along with instructions how to reproduce can be found in the [examples/benchmark_c_py](#) directory in the PyBayes source code repository.

```
----- integrate_python.py -----
def f(x):
    return x*x

def integrate(a, b, N):
    s = 0
    dx = (b-a)/N
    for i in xrange(N):
        s += f(a + (i + 1./2.)*dx)*dx
    return s
```

The test was performed on a 64-bit dual-core Intel Core i5-2520M CPU clocked at 2.50 Ghz with Intel Turbo Boost and Hyper-threading enabled, giving the total of 4 logical processor cores (each physical core being able to execute 2 CPU threads); operating system is Gentoo Linux compiled for the x86_64 platform. Versions of relevant software packages are listed below:

Python 2.7.1

GNU C Compiler 4.4.5; -O2 optimisation flag used when compiling C files

Cython 0.14.1+ (git revision 0.14.1-1002-g53e4c10)

PyPy 1.5.0-alpha0 with JIT compiler enabled

Program abbreviations used in benchmark results:

c_omp	C version, for loop parallelised using OpenMP
cy_typed_omp	Cython version with all variables typed, parallelised
pypy	PyPy-executed Python version, single-threaded
c	C version, single-threaded
cython_typed	Cython version with all variables typed, single-threaded
cython	Cython-compiled Python version, single-threaded
python	Python version, single-threaded

```
----- typical benchmark run -----
Numerical integration from 0.0 to 3.0 of x^2 with 200000000 steps:

    c_omp: result = 9.0; real time = 0.446s; cpu time = 1.72s
  cy_typed_omp: result = 9.0; real time = 0.447s; cpu time = 1.73s
      pypy: result = 9.0; real time = 0.851s; cpu time = 0.84s
        c: result = 9.0; real time = 1.597s; cpu time = 2.00s
  cython_typed: result = 9.0; real time = 1.590s; cpu time = 1.58s
```

cython: result = 9.0; real time = 33.26s; cpu time = 33.2s
python: result = 9.0; real time = 95.05s; cpu time = 94.8s
Relative speedups:
cython/python: 2.8571203030
cython_typed/cython: 20.9136660253
c/cython_typed: 0.99693008347
cy_typed_omp/cython_typed: 3.55556994329
c_omp/c: 3.57674225915
c_omp/cy_typed_omp: 1.00186053812

The `cpu time` quantity is not reliable in our belief and is mentioned only for illustration; all measurements are based on the `real time` quantity that measures wall-clock time needed to perform the algorithm. The benchmark had little variance in results across runs, the relative sample standard deviation (2.1) was under 3% for all measured quantities (run times) with $N = 10$.

$$s_{\text{rel}} = \frac{1}{\bar{x}} \sqrt{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2} \quad (2.1)$$

The test produced a couple of interesting results: the same Python code was $2.9\times$ faster when compiled by Cython; we are convinced that this is due to the Python interpreter overhead. Adding static type declarations to the Cython code resulted in additional 21-fold speed-up giving the total $60\times$ increase in performance of a statically-typed Cython procedure compared to the dynamically typed Python one, forming a very impressive result. Somewhat surprising results were obtained from parallelisation tests — parallelised Cython version was $3.5\times$ faster than equivalent single-threaded code on a system with only 2 physical processor cores, we speculate that this is due to the Hyper-threading technology employed by the processor that reduced wasted CPU cycles (where the CPU waits for memory fetches). Nonetheless this shows that Cython (and C) parallelisation techniques using OpenMP scale very well. C and Cython versions of the algorithm performed virtually equally in both single-threaded and multi-threaded cases giving an indication that Cython performs a very good job at optimising Python code in such simple cases. Another surprise was PyPy (operating on unmodified Python code) performance — it was considerably faster than both C and optimised Cython code making it a very interesting option once it supports NumPy.

We should however note that these extreme Cython performance gains are specific for such simple and highly numeric algorithms. We expect smaller benefits for more high-level code; Kalman filter tests (section 3.4 on page 33) support our claims.

2.6.3 Discussion

Cython, in our view, fixes the last possible barrier (the CPython overhead) before Python can be used for the desired library for Bayesian filtering. We consider especially important that optimisation can be approached *gradually* — one can write code purely in Python and add static type definitions only for performance critical parts (that show up as bottlenecks during profiling) and only once the high performance is needed; Cython developers soundly discourage adding static types everywhere. Second key feature of Cython, as identified by us and for our purposes, is the pure Python mode. Employing Cython brings the disadvantages of compiled languages, mainly a prolonged write-build-run-debug cycle. When the pure Python mode⁵⁷ is used, all these shortcomings are effectively voided (as the code works also under plain Python) at the cost of very minor loss of programming convenience. We expect that many potential library users don't have high performance requirements, these could ignore Cython entirely and use the library as any other plain python module.

Python/Cython combination was therefore chosen as the implementation environment for the desired library for recursive Bayesian estimation; the library was named *PyBayes* and is presented in the next chapter.

⁵⁷specifically, the variant where only augmenting files are used

Chapter 3

The PyBayes Library

Introduction, general directions, future considerations
+ open development on github, open-source

3.1 Interpreted and Compiled

3.2 Library Layout

[proposed citation: [\[18\]](#)]

3.2.1 Random Variable Meta-representation

Why it is needed (ref to ProdCPdf)

3.2.2 Probability Density Functions

Nice UML diagrams! (better more smaller UMLs than one big) One for general pdf layout, one for AbstractGaussPdf family, one for AbstractEmpPdf family

3.2.3 Bayesian Filters

UML

Nice graph of a run of a particle filter (Mirda has the plotting code)
similar of marginalized particle filter? (gausses would be plotted vertically)
[mention this:[\[15\]](#)]

3.3 Documentation, Testing and Profiling

TODO: move above Library Layout?

Documenting PyBayes using Sphinx, approach to documentation (mathematician-oriented), math in documentation

Testing - the separation of

- tests: test one class in isolation, quick, determinism (would be good, not achievable)

- stresses: test a great portion of code at once, run longer, non-determinism..

Note about coverage.py!!

Profiling python/cython - how, existing support in PyBayes

- how to correct profiling-induced overhead

3.4 Performance Comparison with BDM

[skip if in time press]

Conclusion

Bibliography

- [1] M. Sanjeev Arulampalam, Simon Maskell, and Neil Gordon. A tutorial on particle filters for online nonlinear/non-gaussian bayesian tracking. *IEEE Transactions on Signal Processing*, 50:174–188, 2002. 6, 7, 8, 9, 10, 11
- [2] S. Behnel, R. Bradshaw, C. Citro, L. Dalcin, D.S. Seljebotn, and K. Smith. Cython: The best of both worlds. *Computing in Science Engineering*, 13(2):31–39, march-april 2011. 25
- [3] Stefan Behnel, Robert W. Bradshaw, and Dag Sverre Seljebotn. Cython tutorial. In Gaël Varoquaux, Stéfan van der Walt, and Jarrod Millman, editors, *Proceedings of the 8th Python in Science Conference*, pages 4–14, Pasadena, CA USA, 2009. 27
- [4] D. Crisan and A. Doucet. A survey of convergence results on particle filtering methods for practitioners. *Signal Processing, IEEE Transactions on*, 50(3):736–746, 2002. 10
- [5] K. Gadeyne. BFL: Bayesian Filtering Library. <http://www.orocos.org/bfl>, 2001. 21
- [6] F. Gustafsson, F. Gunnarsson, N. Bergman, U. Forssell, J. Jansson, R. Karlsson, and P.J. Nordlund. Particle filters for positioning, navigation, and tracking. *Signal Processing, IEEE Transactions on*, 50(2):425–437, 2002. 4
- [7] R. Hofman, V. Šmídl, and P. Pecha. Data assimilation in early phase of radiation accident using particle filter. In *The Fifth WMO International Symposium on Data Assimilation*, Melbourne, Australia, 2009. 4
- [8] R. Hofman and Šmídl V. Assimilation of spatio-temporal distribution of radionuclides in early phase of radiation accident. *Bezpečnost jaderné energie*, 18:226–228, 2010. 4
- [9] Y. Levanoni and E. Petrank. An on-the-fly reference counting garbage collector for Java. *ACM Transactions on Programming Languages and Systems*, 28(1), January 2006. 19
- [10] P. Pecha, Hofman R., and V. Šmídl. Bayesian tracking of the toxic plume spreading in the early stage of radiation accident. In *Proceedings of the 2009 European Simulation and Modelling Conference*, Leicester, GB, 2009. 4

- [11] V. Peterka. Bayesian approach to system identification. In P. Eykhoff, editor, *Trends and Progress in System identification*, pages 239–304. Pergamon Press, Oxford, 1981. 8
- [12] T. B. Schön, F. Gustafsson, and P.-J. Nordlund. Marginalized particle filters for mixed linear/nonlinear state-space models. *Signal Processing, IEEE Transactions on*, 53(7):2279–2289, july 2005. 11
- [13] T. B. Schön, R. Karlsson, and F. Gustafsson. The marginalized particle filter — analysis, applications and generalizations, 2006. 11
- [14] Dag Sverre Seljebotn. Fast numerical computations with cython. In Gaël Varoquaux, Stéfan van der Walt, and Jarrod Millman, editors, *Proceedings of the 8th Python in Science Conference*, pages 15–22, Pasadena, CA USA, 2009. 27
- [15] V. Šmídl. Software analysis unifying particle filtering and marginalized particle filtering. In *Proceedings of the 13th International Conference on Information Fusion*. IET, 2010. 7, 32
- [16] B. Stroustrup. *The C++ programming language, Third Edition*. Addison-Wesley, 2000. 20
- [17] S. Thrun, W. Burgard, and D. Fox. *Probabilistic Robotics (Intelligent Robotics and Autonomous Agents)*. 2005. 4
- [18] V. Šmídl. *Software analysis of Bayesian distributed dynamic decision making*. PhD thesis, University of West Bohemia, Faculty of Applied Sciences, Pilsen, Czech Republic, Plzeň, 2005. 32
- [19] V. Šmídl and M. Pištěk. Presentation of Bayesian Decision Making Toolbox (BDM). In M. Janžura and J. Ivánek, editors, *Abstracts of Contributions to 5th International Workshop on Data — Algorithms — Decision Making*, page 37, Praha, 2009. 21, 23